

# OpenMP Overview

in 30 Minutes

Christian Terboven <terboven@rz.rwth-aachen.de>

06.12.2010 / Aachen, Germany

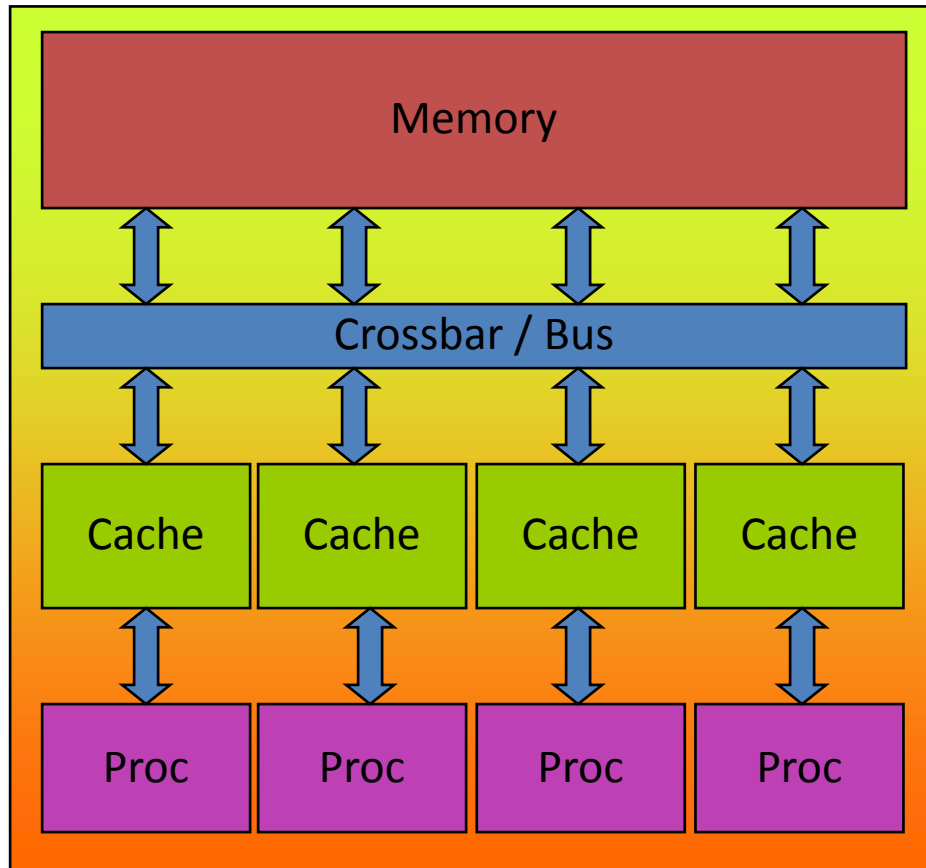
Stand: 03.12.2010

Version 2.3

- ▶ **OpenMP: Parallel Regions, Worksharing, Synchronization**
- ▶ **Example: Pi**
- ▶ **OpenMP: Tasking**
- ▶ **Example: Fibonacci**

# OpenMP: Parallel Regions, Worksharing, Synchronization


## ► OpenMP: Shared-Memory Parallel Programming Model.



All processors/cores access a shared main memory.

Real architectures are more complex, as we will see later / as you just have seen.

Parallelization in OpenMP employs threads.



- ▶ OpenMP programs start with just one thread: The *Master*.
- ▶ *Worker* threads are spawned at *Parallel Regions*. Together with the Master they form a *Team*.
- ▶ In between Parallel Regions the Worker threads are put to sleep.
- ▶ Concept: *Fork-Join*.
- ▶ Allows for an incremental parallelization!

- ▶ **The parallelism has to be expressed explicitly.**

C/C++

```
#pragma omp parallel
{
    ...
    structured block
    ...
}
```

- ▶ ***Structured Block***

- ▶ Exactly one entry point at the top
- ▶ Exactly one exit point at the bottom
- ▶ Branching in or out is not allowed
- ▶ Terminating the program is allowed (abort)

- ▶ **Specification of number of threads:**

- ▶ Environment variable:

OMP\_NUM\_THREADS=...

- ▶ Or: Via `num_threads` clause:

```
#pragma omp parallel \
    num_threads(num) {...}
```

- ▶ If only the *parallel* construct is used, each thread executes the Structured Block.
- ▶ Program Speedup: *Worksharing*
- ▶ OpenMP's most common Worksharing construct: *for*

C/C++

```
int i;  
#pragma omp parallel for  
for (i = 0; i < 100; i++)  
{  
    a[i] = b[i] + c[i];  
}
```

- ▶ Distribution of loop iterations over all threads in a Team.
- ▶ Scheduling of the distribution can be influenced.
- ▶ Loops often account for most of the program runtime!

- ▶ **for-construct: OpenMP allows to influence how the iterations are scheduled among the threads of the team, via the *schedule* clause:**
  - ▶ `schedule(static [, chunk])`: Iteration space divided into blocks of chunk size, blocks are assigned to threads in a round-robin fashion. If chunk is not specified: #threads blocks.
  - ▶ `schedule(dynamic [, chunk])`: Iteration space divided into blocks of chunk (not specified: 1) size, blocks are scheduled to threads in the order in which threads finish previous blocks.
  - ▶ `schedule(guided [, chunk])`: Similar to dynamic, but block size starts with implementation-defined value, then is decreased exponentially down to chunk.
- ▶ **Default on most implementations is `schedule(static)`.**



- ▶ **Challenge of Shared-Memory parallelization: Managing the Data Environment.**
- ▶ ***Scoping in OpenMP: Dividing variables in *shared* and *private*:***
  - ▶ *private*-list and *shared*-list on Parallel Region
  - ▶ *private*-list and *shared*-list on Worksharing constructs
  - ▶ Default is *shared*
  - ▶ Loop control variables on *for*-constructs are *private*
  - ▶ Non-static variables local to Parallel Regions are *private*
  - ▶ *private*: A new uninitialized instance is created for each thread
    - ▶ *firstprivate*: Initialization with Master's value
    - ▶ *lastprivate*: Value of last loop iteration is written back to Master
  - ▶ Static variables are *shared*

- ▶ **Global / static variables can be privatized with the *threadprivate* directive**
  - ▶ One instance is created for each thread
    - ▶ Before the first parallel region is encountered
    - ▶ Instance exists until the program ends
    - ▶ Does not work (well) with nested Parallel Region
  - ▶ Based on thread-local storage (TLS)
    - ▶ TlsAlloc (Win32-Threads), pthread\_key\_create (Posix-Threads), keyword `__thread` (GNU extension)

C/C++

```
static int i;  
#pragma omp threadprivate(i)
```

## ► Can all loops be parallelized with `for`-constructs? No!

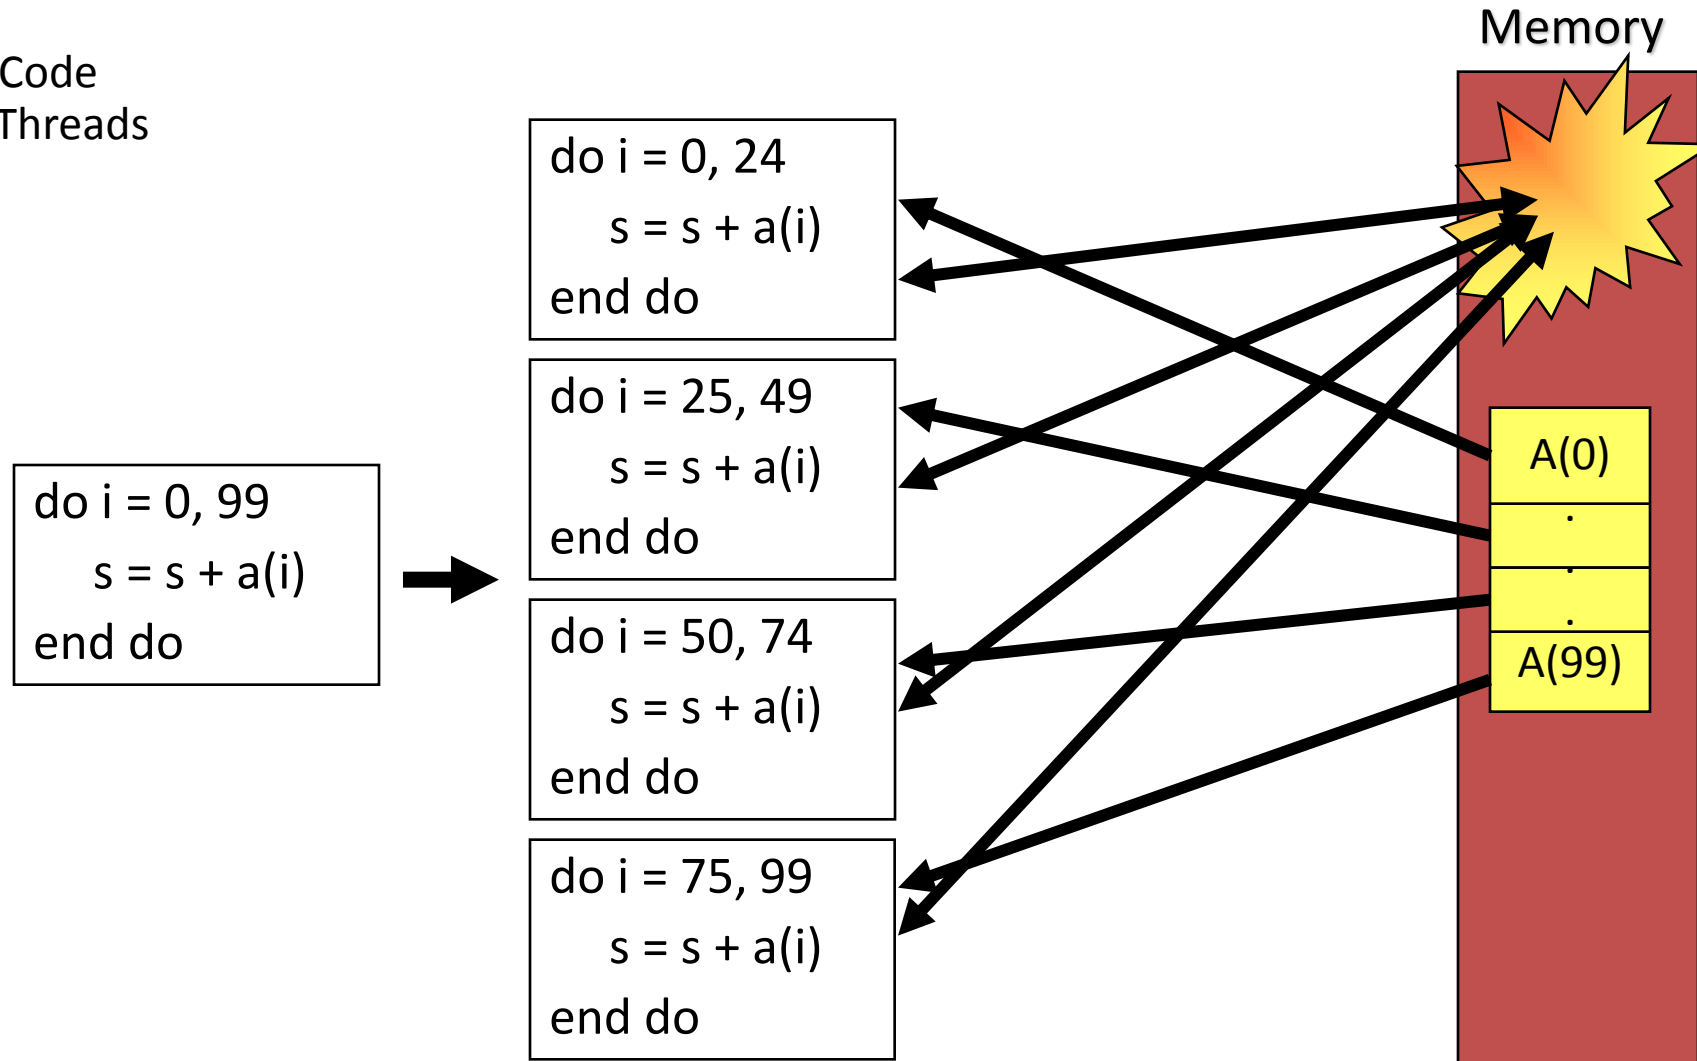
- Simple test: If the results differ when the code is executed backwards, the loop iterations are not independent. BUT: This test alone is not sufficient:

C/C++

```
int i;  
#pragma omp parallel for  
for (i = 0; i < 100; i++)  
{  
    s = s + a[i];  
}
```

- **Data Race:** If between two synchronization points at least one thread writes to a memory location from which at least one other thread reads, the result is not deterministic (race condition).

Pseudo-Code  
Here: 4 Threads



- ▶ **A *Critical Region* is executed by all threads, but by only one thread simultaneously (*Mutual Exclusion*).**

C/C++

```
#pragma omp critical (name)
{
    ... structured block ...
}
```

- ▶ **Do you think this solution scales well?**

C/C++

```
int i;
#pragma omp parallel for
for (i = 0; i < 100; i++)
{
    #pragma omp critical
    { s = s + a[i]; }
}
```

- ▶ In a *reduction*-operation the operator is applied to all variables in the list. The variables have to be *shared*.

- ▶ `reduction(operator:list)`

- ▶ The result is provided in the associated reduction variable

C/C++

```
#pragma omp parallel for reduction(+:s)
for(i = 0; i < 99; i++)
{
    s = s + a[i];
}
```

- ▶ Possible reduction operators: `*`, `-`, `&`, `|`, `&&`, `||`, `^`

## ► C and C++:

- If OpenMP is enabled during compilation, the preprocessor symbol `_OPENMP` is defined. To use the OpenMP runtime library, the header `omp.h` has to be included.
- `omp_set_num_threads(int)`: The specified number of threads will be used for the parallel region encountered next.
- `int omp_get_num_threads`: Returns the number of threads in the current team.
- `int omp_get_thread_num()`: Returns the number of the calling thread in the team, the Master has always the id 0.

## ► Additional functions are available, e.g. to provide locking functionality.

# Example: Pi



### ○ Simple example: calculate Pi by integration

```
double f(double x) {  
    return (double)4.0 / ((double)1.0 + (x*x));  
}  
  
void computePi() {  
    double h = (double)1.0 / (double)iNumIntervals;  
    double sum = 0, x;  
  
    for (int i = 1; i <= iNumIntervals; i++) {  
        x = h * ((double)i - (double)0.5);  
        sum += f(x);  
    }  
  
    myPi = h * sum;  
}
```

$$\Pi = \int_0^1 \frac{4}{(1+x^2)} dx$$

# OpenMP: Tasking

# How to parallelize a While-loop?

## ► How would you parallelize this code?

```
typedef list<double> dList;  
dList myList;  
/* fill myList with tons of items */  
  
dList::iterator it = myList.begin();  
while (it != myList.end())  
{  
    *it = processListItem(*it);  
    it++;  
}
```

- **One possibility: Create a fixed-sized array containing all list items and a parallel loop running over this array**  
**Concept: Inspector / Executor**

# How to parallelize a While-loop!

## ► Or: Use Tasking in OpenMP 3.0

```
#pragma omp parallel
{
    #pragma omp single
    {
        dList::iterator it = myList.begin();
        while (it != myList.end())
        {
            #pragma omp task
            {
                *it = processListItem(*it);
            }
            it++;
        }
    }
}
```

## ► All while-loop iterations are independent from each other!

# The task directive

C/C++

```
#pragma omp task [clause [[,] clause] ... ]  
... structured block ...
```

## ▶ Each encountering thread creates a new Task

- ▶ Code and data is being packaged up
- ▶ Tasks can be nested
  - ▶ Into another Task directive
  - ▶ Into a Worksharing construct

## ▶ Data scoping clauses:

- ▶ `shared(list)`
- ▶ `private(list)`
- ▶ `firstprivate(list)`
- ▶ `default(shared | none)`

## ▶ At OpenMP `barrier` (implicit or explicit)

- ▶ All tasks created by any thread of the current *Team* are guaranteed to be completed at barrier exit

## ▶ Task barrier: `taskwait`

- ▶ Encountering Task suspends until child tasks are complete
  - ▶ Only direct childs, not descendants!

C/C++

```
#pragma omp taskwait
```

### ► Simple example of Task synchronization in OpenMP 3.0:

```
#pragma omp parallel num_threads(np)
{
    #pragma omp task
        function_A();
    #pragma omp barrier
    #pragma omp single
    {
        #pragma omp task
            function_B();
    }
}
```

np Tasks created here, one for each thread

All Tasks guaranteed to be completed here

1 Task created here

B-Task guaranteed to be completed here

- ▶ **Some rules from *Parallel Regions* apply:**
  - ▶ Static and Global variables are shared
  - ▶ Automatic Storage (local) variables are private
  
- ▶ **If no default clause is given:**
  - ▶ Orphaned Task variables are `firstprivate` by default!
  - ▶ Non-Orphaned Task variables inherit the `shared` attribute!

→ Variables are `firstprivate` unless `shared` in the enclosing context
  
- ▶ **So far no verification tool is available to check Tasking programs for correctness!**



# Example: Fibonacci

```
int main(int argc,  
        char* argv[])  
{  
    [...]  
    fib(input);  
    [...]  
}
```

```
int fib(int n)    {  
    if (n < 2) return n;  
    int x = fib(n - 1);  
    int y = fib(n - 2);  
    return x+y;  
}
```


- ▶ **On the following slides we will discuss three approaches to parallelize this recursive code with Tasking.**

```
int main(int argc,
         char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)    {
    if (n < 2) return n;
    int x, y;
    #pragma omp task shared(x)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

- Only one Task / Thread enters `fib()` from `main()`, it is responsible for creating the two initial work tasks
- Taskwait is required, as otherwise `x` and `y` would be lost

- **Overhead of task creation prevents better scalability!**




- **Improvement: Don't create yet another task once a certain (small enough)  $n$  is reached**

```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)    {
    if (n < 2) return n;
    int x, y;
    #pragma omp task shared(x) \
        if(n > 30)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y) \
        if(n > 30)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

- ▶ Speedup is ok, but we still have some overhead when running with 4 or 8 threads




- **Improvement: Skip the OpenMP overhead once a certain  $n$  is reached (no issue w/ production compilers)**

```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)    {
    if (n < 2) return n;
    if (n <= 30)
        return serfib(n);
    int x, y;
    #pragma omp task shared(x)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

► Everything ok now 😊

Speedup of Fibonacci with Tasks





# The End

**Thank you for your attention.**